

1 Preliminaries

Before starting on this assignment, please be sure to read the General Instructions that are on Piazza (under Resources->General Resources). If you did Lab1, you should already know how to log in to the class PostgreSQL server. You'll get help on Lab2 in your Discussion Section, not the Lectures, so *be sure to attend Discussion Sections*.

2 Goal

The goal of the second assignment is to create a PostgreSQL data schema with 6 tables that are very similar to the tables that you created in Lab1. The tables have the same names, attributes and data types as the tables of Lab1, and the same primary keys but there are **some UNIQUE constraints and some restrictions on NULLs**.

After you create the data schema with the 6 tables, you will be required to write some SQL statements that use those tables. We will give you data that you can load into your tables so that you can test the results of your queries. Testing can prove that a query is wrong, but not that it is right, so be careful.

Lab2 is due in two weeks, so you will have an opportunity to discuss the assignment during the Discussion Section in the first week of the assignment, and to discuss issues you have had in writing a solution to the assignment during the Discussion Section of the second week. Instructions for submitting the assignment appear at the end of this document.

3 Lab2 Description

3.1 Create PostgreSQL Schema Lab2

You will create a Lab2 schema to set apart the database tables created in this lab from ones you will create in future, as well as from tables (and other objects) in the default (public) schema. Note that the meaning of schema here is specific to PostgreSQL and different from the general meaning. See [here](#) for more details on PostgreSQL schemas. You create the Lab2 schema like this:

```
CREATE SCHEMA Lab2;
```

Now that you have created the schema, you want to set Lab2 to be your default schema when you use psql. If you do not set Lab2 as the default schema, then you will have to qualify your table names with the schema name (e.g. Lab2.Customers). To set the default schema, you modify your search path. (For more details, see [here](#).)

```
ALTER ROLE username SET SEARCH_PATH to Lab2;
```

You do not have to include the CREATE SCHEMA or ALTER ROLE statements in your solution.

3.2 Create tables

You will create tables in schema Lab2 for the tables Authors, Books, Publishers, Members, Orders and Reviews. The attributes of the 6 tables are the same as the tables of Lab1. Data types for the attribute names in these tables are also the same as the ones specified for the tables of Lab1. The Primary Keys are also the same. However, the tables must have the additional constraints described in the next section.

3.2.1 Constraints

The following attributes cannot be NULL. All other attributes can be (but remember that attributes in Primary Keys also cannot be NULL).

- In Books: publisherID
- In Reviews: reviewStars
- In Orders: quantity

Also, the following must be unique for the specified table. That is, there cannot be identical rows in that table that have exactly the same (non-NULL) values for all of those attributes (composite unique constraint).

- In Publishers: the attribute publisherName
- In Books: the 2 attributes authorID and bookName
- In Members: the 2 attributes memberName and joinDate

For example, the third constraint says that there can't be two rows in Members that have the same values for both memberName and joinDate (if both memberName and joinData are not NULL). Think of this as saying that there can't be two members with the same name who joined on the same date.

You will write a CREATE TABLE command for each of the 6 tables. Save the commands in the file create.sql

4 SQL Queries

Below are English descriptions of the five SQL queries that you need to write for this assignment, which you will include in `files queryX.sql`, where X is the number of the query, e.g., your SQL statement for Query 1 will be in the file `query1.sql` and so forth. Follow the directions as given. You will lose points if you give extra tuples or attributes in your results, if you give attributes in with the wrong names or in the wrong order, or if you have missing or wrong results. You will also lose points if your queries are unnecessarily complex, even if they are correct.

Note: An author is someone who appears as a row in Authors, a book appears as a row in Books, a member is someone who appears as a row in Members, etc.

4.1 Query 1

Find the authorID and authorName for each author whose name doesn't begin with the letter S. Your result should appear in descending order of authorName, and no duplicates should appear in your result.

4.2 Query 2

Output the quantity and orderDate for each order that was made by a member whose isCurrentMember is TRUE. The second attribute in your result should appear as oDate. Duplicates may appear in your result.

4.3 Query 3

Category is 'N' for non-fiction books and 'F' for fiction books. Find the authorName and publisherName for each non-fiction book whose price is greater than or equal to 19.99; no duplicates should appear in your result.

4.4 Query 4

Find the reviewerID and bookID for each review that has all of the following three properties:

- The number of reviewStars isn't 4.
- The book is named 'Jane Eyre'.
- The reviewer is a member whose name is 'John Smith'.

No duplicates should appear in your result.

4.5 Query 5

Find the memberID and joinDate for each member who reviewed a book that was published after April 1, 2010, and that book was also written by an author who has the exact same name as the reviewer. No duplicates should appear in your result.

5 Testing

While your solution is still a work in progress, it is a good idea to drop all objects from the database every time you run the script, so you can start fresh. Of course, dropping each object may be tedious, and sometimes there may be a particular order in which objects must be dropped. The following commands (which you can put at the top of `create.sql` if you want, but you don't have to), will drop your Lab2 schema (and all objects within it), and then create the (empty) schema again:

```
DROP SCHEMA Lab2 CASCADE;  
CREATE SCHEMA Lab2;
```

Before you submit, login to your database via `psql` and execute your script. As you've learned already, the command to execute a script is: `\i <filename>`.

Under Resources→Lab2 on Piazza, we will also provide (in the next day or two) a load script named *lab2_data_loading.sql*, that loads data into the 6 tables of the database. You can execute that script with the command:

```
\i lab2_data_loading.sql.
```

You can test your 5 queries using that data, but you will have to figure out on your own whether your query results are correct. We won't provide answers, and students should not share answers with other students. Also, your queries must be correct on any database instance, not just on the data that we provide. You may want to test your SQL statements on your own data as well.

6 Submitting

1. Save your scripts for table creations and query statements as create.sql and query1.sql through query5.sql. You may add informative comments inside your scripts if you want (the server interprets lines that start with two hyphens as comment lines).
2. Zip the file(s) to a single file with name Lab2_XXXXXXX.zip where XXXXXXX is your 7-digit student ID. For example, if a student's ID is 1234567, then the file that this student submits for Lab2 should be named Lab2_1234567.zip

To generate the zip file you can use the Unix command:

```
zip Lab2_1234567 create.sql query1.sql query2.sql query3.sql query4.sql query5.sql
```

(Of course, you use your own student ID, not 1234567.)

3. You should already know how to transfer the files from the UNIX timeshare to your local machine before submitting to Canvas. If you are still not familiar with the process, use the instructions we provided at the Lab1 assignment.
4. Lab2 is due by 11:59pm on Sunday, October 28. Late submissions will not be accepted, and there will be no make-up Lab assignments.
5. You may lose credit if your queries are unnecessarily complex, even if they are correct. For example, if you use DISTINCT but it's not needed, you might lose half a point. Of course, you will also lose credit if DISTINCT is needed and you omit it.
6. Be sure to follow directions about Academic Integrity that are in the Syllabus and Lecture1. If you have any questions about those directions, please speak to the instructor as soon as possible.